

MRF Compiler Reference Manual

Shyamal Suhana Chandra

November 17, 2025

Contents

1	Introduction	3
1.1	Purpose	3
1.2	Scope	3
2	Installation	3
2.1	System Requirements	3
2.2	Building from Source	4
2.2.1	Linux/macOS	4
2.2.2	Windows	4
2.3	Verification	4
3	Command-Line Interface	4
3.1	Basic Usage	4
3.2	Command-Line Options	4
3.3	Examples	5
3.3.1	Basic Usage	5
3.3.2	Export to All Frameworks	5
3.3.3	Verbose Mode	5
4	Input Format	5
4.1	File Structure	5
4.2	Graph Type Declaration	6
4.3	Node Declaration	6
4.4	Edge Declaration	6
4.5	Complete Example	6
4.6	Directed Graph Example	7
5	Output Formats	7
5.1	OpenQASM 2.0	7
5.2	Qiskit	7
5.3	Cirq	8
5.4	PennyLane	8
5.5	Q#	9
5.6	AWS Braket	9
5.7	Qulacs	9
5.8	TensorFlow Quantum	10

6 API Reference	10
6.1 Graph Module	10
6.1.1 GraphicalModel Class	10
6.2 MRF Module	11
6.2.1 MRF Class	11
6.2.2 Conversion Functions	11
7 Framework Integration	11
7.1 Qiskit Integration	11
7.2 Cirq Integration	12
7.3 PennyLane Integration	12
8 Examples	12
8.1 Example 1: Simple Chain Graph	12
8.2 Example 2: Directed Graph with Moralization	12
8.3 Example 3: Export to All Frameworks	13
9 Troubleshooting	13
9.1 Common Issues	13
9.1.1 Compilation Errors	13
9.1.2 File Not Found	13
9.1.3 Invalid Input Format	13
9.1.4 Framework Import Errors	13
9.1.5 Large Graph Performance	14
10 Advanced Usage	14
10.1 Custom Potential Functions	14
10.2 Circuit Optimization	14
10.3 Performance Considerations	14
11 Best Practices	14
12 Contributing	15
13 License	15
14 References	15

1 Introduction

This reference manual provides complete documentation for the MRF Compiler, a comprehensive tool for converting probabilistic graphical models into Markov Random Fields (MRF) and subsequently into quantum circuit representations. The compiler supports multiple quantum computing frameworks and provides a flexible command-line interface for easy integration into existing workflows.

1.1 Purpose

The MRF Compiler bridges the gap between classical probabilistic modeling and quantum computing, enabling researchers and practitioners to:

- Convert graphical models to quantum circuits automatically
- Export to multiple quantum computing frameworks
- Leverage quantum algorithms for probabilistic inference
- Integrate quantum computing into existing machine learning pipelines

1.2 Scope

This manual covers:

- Installation and setup procedures
- Command-line interface and options
- Input file format specifications
- Output format details for each framework
- API reference for programmatic usage
- Framework-specific integration guides
- Examples and use cases
- Troubleshooting and common issues

2 Installation

2.1 System Requirements

- **Operating System:** Linux, macOS, or Windows (with WSL or MinGW)
- **Compiler:** C++17 compatible compiler (GCC 7+, Clang 5+, or MSVC 2017+)
- **Build System:** Make
- **Memory:** Minimum 512MB RAM (2GB recommended for large graphs)
- **Disk Space:** 50MB for installation

2.2 Building from Source

2.2.1 Linux/macOS

```
1 # Clone the repository
2 git clone https://github.com/shyamalchandra/MRFcompiler.git
3 cd MRFcompiler
4
5 # Build the compiler
6 make
7
8 # Optional: Install system-wide
9 sudo make install
```

2.2.2 macOS

The MRF Compiler supports macOS with both Homebrew LLVM (clang++) and GCC, as well as Apple's system clang.

Option 1: Using Homebrew LLVM (Recommended)

```
1 # Install Homebrew (if not already installed)
2 /bin/bash -c "$(curl -fsSL
3   https://raw.githubusercontent.com/Homebrew/install/HEAD/install.sh)"
4
5 # Install LLVM
6 brew install llvm
7
8 # Clone and build
9 git clone
10  https://github.com/Sapana-Micro-Software/MRFcompiler.git
11 cd MRFcompiler
12 make check-compiler # Verify compiler detection
13 make
```

Option 2: Using Homebrew GCC

```
1 # Install GCC via Homebrew
2 brew install gcc
3
4 # Clone and build
5 git clone
6   https://github.com/Sapana-Micro-Software/MRFcompiler.git
7 cd MRFcompiler
8 make
```

Option 3: Using Xcode Command Line Tools

```
1 # Install Xcode Command Line Tools
2 xcode-select --install
3
4 # Clone and build
5 git clone
6   https://github.com/Sapana-Micro-Software/MRFcompiler.git
```

```
6 cd MRFcompiler  
7 make
```

The Makefile automatically detects and uses the best available compiler:

- Homebrew LLVM (if installed at `/opt/homebrew/opt/llvm` or `/usr/local/opt/llvm`)
- Homebrew GCC (if installed, checks for `g++-13`, `g++-12`, `g++-11`)
- System clang (Apple's LLVM) as fallback

2.2.3 Windows

```
1 # Using MinGW or MSYS2  
2 git clone  
   https://github.com/Sapana-Micro-Software/MRFcompiler.git  
3 cd MRFcompiler  
4 mingw32-make  
5  
6 # Or using Visual Studio  
7 # Open the project in Visual Studio and build
```

2.3 Verification

After building, verify the installation:

```
1 ./mrf_compiler --help
```

You should see the help message with available options.

3 Command-Line Interface

3.1 Basic Usage

The basic command syntax is:

```
./mrf_compiler [options] [input_file] [output_file]
```

3.2 Command-Line Options

Option	Description
<code>-f</code> , <code>--framework</code>	Specify output framework. Supported values: <code>qasm</code> , <code>qiskit</code> , <code>cirq</code> , <code>pennylane</code> , <code>qsharp</code> , <code>braket</code> , <code>qulacs</code> , <code>tfq</code> . Default: <code>qasm</code>
<code>-a</code> , <code>--all</code>	Export to all supported frameworks. Creates multiple output files with framework-specific extensions.
<code>-h</code> , <code>--help</code>	Display help message with usage information and available options.

Option	Description
<code>-v, --verbose</code>	Enable verbose output, showing detailed conversion steps and debugging information.
<code>-o, --output <file></code>	Specify output file path. If not provided, output is written to stdout or a default filename based on input.

3.3 Examples

3.3.1 Basic Usage

```

1 # Convert to OpenQASM (default)
2 ./mrf_compiler example.txt output.qasm
3
4 # Convert to Qiskit
5 ./mrf_compiler -f qiskit example.txt circuit.py
6
7 # Convert to Cirq
8 ./mrf_compiler -f cirq example.txt circuit.py

```

3.3.2 Export to All Frameworks

```

1 # Generate output for all frameworks
2 ./mrf_compiler -a example.txt
3
4 # This creates:
5 # - output_qasm.qasm
6 # - output_qiskit.py
7 # - output_cirq.py
8 # - output_pennylane.py
9 # - output_qsharp.qs
10 # - output_braket.py
11 # - output_qulacs.py
12 # - output_tfq.py

```

3.3.3 Verbose Mode

```

1 # Get detailed conversion information
2 ./mrf_compiler -v -f qiskit example.txt circuit.py

```

4 Input Format

4.1 File Structure

The input file is a plain text file with a simple line-based format. Each line specifies either the graph type, a node, or an edge.

4.2 Graph Type Declaration

The first line must specify the graph type:

TYPE directed

or

TYPE undirected

4.3 Node Declaration

Nodes are declared with the following syntax:

NODE <id> <name> [num_states]

- <id>: Unique integer identifier for the node (starting from 0)
- <name>: Alphanumeric name for the node
- [num_states] : *Optional number of states (default : 2 for binary)*

4.4 Edge Declaration

Edges are declared with the following syntax:

EDGE <from> <to> [directed]

- <from>: Source node ID
- <to>: Target node ID
- [directed]: Optional flag (only relevant for undirected graphs, ignored for directed)

4.5 Conditional Probability Tables (CPTs)

For Bayesian Networks (directed graphs), you can specify Conditional Probability Tables (CPTs) for each node. CPTs define the probability distribution of a node given its parents.

4.5.1 CPT Syntax

The CPT command syntax is:

CPT <node_id> [parent_states...] <prob_state0> <prob_state1> ...

- <node_id>: The node ID for which the CPT is defined
- [parent_sstates...]: For nodes with parents, specify the parent state combination (one integer per parent)
- <prob_state0><prob_state1> ... : Probability values for each state of the node (must sum to 1.0 for each state)

4.5.2 Root Node CPTs

For root nodes (nodes with no parents), simply specify the probability distribution:

```
CPT 0 0.8 0.2
```

This means: $P(\text{Node0}=0) = 0.8$, $P(\text{Node0}=1) = 0.2$

4.5.3 Node with Parents CPTs

For nodes with parents, you must specify a CPT entry for each parent state combination.
Each line specifies one parent combination:

```
CPT 2 0 0 0.99 0.01      # P(Node2=0|Parent1=0,Parent2=0)=0.99, P(Node2=1|...)=0.01
CPT 2 0 1 0.9 0.1        # P(Node2=0|Parent1=0,Parent2=1)=0.9, P(Node2=1|...)=0.1
CPT 2 1 0 0.8 0.2        # P(Node2=0|Parent1=1,Parent2=0)=0.8, P(Node2=1|...)=0.2
CPT 2 1 1 0.0 1.0        # P(Node2=0|Parent1=1,Parent2=1)=0.0, P(Node2=1|...)=1.0
```

The parent states are specified in the order of parent node IDs (as they appear in EDGE declarations).

4.5.4 CPT Example: Bayesian Network

A complete example of a Bayesian Network with CPTs:

```
TYPE directed
NODE 0 Rain 2
NODE 1 Sprinkler 2
NODE 2 WetGrass 2
EDGE 0 2 directed
EDGE 1 2 directed
# Root node: P(Rain=0)=0.8, P(Rain=1)=0.2
CPT 0 0.8 0.2
# Root node: P(Sprinkler=0)=0.6, P(Sprinkler=1)=0.4
CPT 1 0.6 0.4
# WetGrass depends on Rain and Sprinkler
CPT 2 0 0 0.99 0.01  # P(WetGrass|Rain=0,Sprinkler=0)
CPT 2 0 1 0.9 0.1    # P(WetGrass|Rain=0,Sprinkler=1)
CPT 2 1 0 0.8 0.2    # P(WetGrass|Rain=1,Sprinkler=0)
CPT 2 1 1 0.0 1.0    # P(WetGrass|Rain=1,Sprinkler=1)
```

4.6 Complete Example

```
TYPE undirected
NODE 0 A 2
NODE 1 B 2
NODE 2 C 2
NODE 3 D 2
EDGE 0 1
EDGE 1 2
EDGE 2 3
EDGE 0 3
```

This creates a 4-node undirected graph forming a square (cycle).

4.7 Directed Graph Example

```
TYPE directed
NODE 0 A 2
NODE 1 B 2
NODE 2 C 2
EDGE 0 1
EDGE 0 2
EDGE 1 2
```

This creates a directed graph where A is a parent of both B and C, and B is a parent of C.

5 Output Formats

5.1 OpenQASM 2.0

OpenQASM is the standard quantum assembly language. The output is a text file that can be used with any OpenQASM-compatible quantum simulator or hardware.

File Extension: .qasm

Example Output:

```
1 OPENQASM 2.0;
2 include "qelib1.inc";
3
4 qreg q[3];
5 creg c[3];
6
7 h q[0];
8 h q[1];
9 h q[2];
10 cx q[0], q[1];
11 rz(0.5) q[1];
12 cx q[0], q[1];
```

5.2 Qiskit

Qiskit is IBM's quantum computing framework. The output is a Python file with a function that returns a `QuantumCircuit` object.

File Extension: .py

Example Output:

```
1 from qiskit import QuantumCircuit
2
3 def create_circuit():
4     qc = QuantumCircuit(3, 3)
5     qc.h(0)
6     qc.h(1)
```

```

7     qc.h(2)
8     qc.cx(0, 1)
9     qc.rz(0.5, 1)
10    qc.cx(0, 1)
11    return qc
12
13 if __name__ == "__main__":
14     circuit = create_circuit()
15     circuit.draw(output='mpl')

```

5.3 Cirq

Cirq is Google's quantum computing framework. The output is a Python file with a function that returns a `cirq.Circuit` object.

File Extension: .py

Example Output:

```

1 import cirq
2
3 def create_circuit():
4     qubits = [cirq.GridQubit(0, i) for i in range(3)]
5     circuit = cirq.Circuit()
6     circuit.append([cirq.H(qubits[i]) for i in range(3)])
7     circuit.append(cirq.CNOT(qubits[0], qubits[1]))
8     circuit.append(cirq.rz(0.5)(qubits[1]))
9     circuit.append(cirq.CNOT(qubits[0], qubits[1]))
10    return circuit

```

5.4 PennyLane

PennyLane is Xanadu's quantum machine learning library. The output is a Python file with a `QNode`-decorated function.

File Extension: .py

Example Output:

```

1 import pennylane as qml
2
3 dev = qml.device('default.qubit', wires=3)
4
5 @qml.qnode(dev)
6 def circuit():
7     qml.Hadamard(wires=0)
8     qml.Hadamard(wires=1)
9     qml.Hadamard(wires=2)
10    qml.CNOT(wires=[0, 1])
11    qml.RZ(0.5, wires=1)
12    qml.CNOT(wires=[0, 1])
13    return qml.state()

```

5.5 Q#

Q# is Microsoft's quantum programming language. The output is a Q# source file with an operation.

File Extension: .qs

Example Output:

```
1 namespace MRFCircuit {
2     open Microsoft.Quantum.Intrinsic;
3     open Microsoft.Quantum.Math;
4
5     operation CreateCircuit() : Unit {
6         using (qubits = Qubit[3]) {
7             H(qubits[0]);
8             H(qubits[1]);
9             H(qubits[2]);
10            CNOT(qubits[0], qubits[1]);
11            Rz(0.5, qubits[1]);
12            CNOT(qubits[0], qubits[1]);
13        }
14    }
15 }
```

5.6 AWS Braket

AWS Braket is Amazon's quantum computing service. The output is a Python file with a function that returns a `braket.Circuit` object.

File Extension: .py

Example Output:

```
1 from braket.circuits import Circuit
2
3 def create_circuit():
4     circuit = Circuit()
5     circuit.h(0)
6     circuit.h(1)
7     circuit.h(2)
8     circuit.cnot(0, 1)
9     circuit.rz(1, 0.5)
10    circuit.cnot(0, 1)
11    return circuit
```

5.7 Qulacs

Qulacs is a fast quantum circuit simulator. The output is a Python file with a function that returns a `qulacs.QuantumCircuit` object.

File Extension: .py

Example Output:

```
1 from qulacs import QuantumCircuit
2
```

```

3 def create_circuit():
4     circuit = QuantumCircuit(3)
5     circuit.add_H_gate(0)
6     circuit.add_H_gate(1)
7     circuit.add_H_gate(2)
8     circuit.add_CNOT_gate(0, 1)
9     circuit.add_RZ_gate(1, 0.5)
10    circuit.add_CNOT_gate(0, 1)
11    return circuit

```

5.8 TensorFlow Quantum

TensorFlow Quantum is Google's quantum machine learning framework. The output is a Python file that creates a tensor representation.

File Extension: .py

Example Output:

```

1 import tensorflow_quantum as tfq
2 import cirq
3
4 def create_circuit():
5     qubits = [cirq.GridQubit(0, i) for i in range(3)]
6     circuit = cirq.Circuit()
7     circuit.append([cirq.H(qubits[i]) for i in range(3)])
8     circuit.append(cirq.CNOT(qubits[0], qubits[1]))
9     circuit.append(cirq.rz(0.5)(qubits[1]))
10    circuit.append(cirq.CNOT(qubits[0], qubits[1]))
11    return tfq.convert_to_tensor([circuit])

```

6 API Reference

6.1 Graph Module

6.1.1 GraphicalModel Class

The GraphicalModel class represents a probabilistic graphical model.

Methods:

- void addNode(int id, const std::string& name, int num_states = 2): Add a node to the graph
- void addEdge(int from, int to, bool directed = true): Add an edge between nodes
- Node* getNode(int id): Get a node by ID
- Edge* getEdge(int from, int to): Get an edge between two nodes
- std::vector<int> getNeighbors(int node_id) const: Get all neighbors of a node

- `bool hasEdge(int from, int to) const`: Check if an edge exists
- `void print() const`: Print the graph structure

6.2 MRF Module

6.2.1 MRF Class

The MRF class represents a Markov Random Field.

Methods:

- `void addNode(int id, const std::string& name, int num_states = 2)`: Add a node
- `void addClique(const std::vector<int>& nodes)`: Add a clique
- `void setCliquePotential(int clique_idx, const std::vector<double>& potential)`: Set clique potential
- `void print() const`: Print the MRF structure
- `int getTotalStates() const`: Get total number of states

6.2.2 Conversion Functions

- `MRF convertToMRF(const GraphicalModel& gm)`: Convert graphical model to MRF
- `void moralizeGraph(GraphicalModel& gm)`: Moralize a directed graph
- `std::vector<Clique> findMaximalCliques(const GraphicalModel& gm)`: Find all maximal cliques

7 Framework Integration

7.1 Qiskit Integration

After generating a Qiskit circuit:

```

1 from output_qiskit import create_circuit
2 from qiskit import Aer, execute
3
4 circuit = create_circuit()
5 backend = Aer.get_backend('qasm_simulator')
6 job = execute(circuit, backend, shots=1000)
7 result = job.result()
8 counts = result.get_counts(circuit)
9 print(counts)

```

7.2 Cirq Integration

After generating a Cirq circuit:

```
1 from output_cirq import create_circuit
2 import cirq
3
4 circuit = create_circuit()
5 simulator = cirq.Simulator()
6 result = simulator.simulate(circuit)
7 print(result)
```

7.3 PennyLane Integration

After generating a PennyLane circuit:

```
1 from output_pennylane import circuit
2 import pennylane as qml
3
4 result = circuit()
5 print(result)
```

8 Examples

8.1 Example 1: Simple Chain Graph

Input (chain.txt):

```
TYPE undirected
NODE 0 A 2
NODE 1 B 2
NODE 2 C 2
EDGE 0 1
EDGE 1 2
```

Command:

```
./mrf_compiler -f qiskit chain.txt chain_circuit.py
```

8.2 Example 2: Directed Graph with Moralization

Input (directed.txt):

```
TYPE directed
NODE 0 A 2
NODE 1 B 2
NODE 2 C 2
EDGE 0 1
EDGE 0 2
```

This creates a v-structure that requires moralization (adding edge between B and C).

Command:

```
./mrf_compiler -f cirq directed.txt directed_circuit.py
```

8.3 Example 3: Export to All Frameworks

Command:

```
./mrf_compiler -a example.txt
```

This generates output files for all 8 supported frameworks.

9 Troubleshooting

9.1 Common Issues

9.1.1 Compilation Errors

Problem: Build fails with compilation errors.

Solution: Ensure you have a C++17 compatible compiler. Update your compiler or use a newer version.

9.1.2 File Not Found

Problem: ./mrf_compiler: No such file or directory

Solution: Ensure you're in the correct directory and the build completed successfully. Run `make` again.

9.1.3 Invalid Input Format

Problem: Error parsing input file.

Solution: Check that your input file follows the correct format:

- First line must be `TYPE directed` or `TYPE undirected`
- Node IDs must be unique integers starting from 0
- Edge node IDs must reference existing nodes

9.1.4 Framework Import Errors

Problem: Generated Python files fail to import framework modules.

Solution: Install the required framework:

```
1 pip install qiskit      # For Qiskit
2 pip install cirq        # For Cirq
3 pip install pennylane   # For PennyLane
4 pip install amazon-braket-sdk # For AWS Braket
5 pip install qulacs      # For Qulacs
6 pip install tensorflow-quantum # For TensorFlow Quantum
```

9.1.5 Large Graph Performance

Problem: Compilation is slow for large graphs.

Solution:

- Clique finding is exponential in worst case
- Consider using smaller graphs or approximate methods
- Increase available memory
- Use verbose mode to identify bottlenecks

10 Advanced Usage

10.1 Custom Potential Functions

While the current version uses default potential functions, future versions will support custom potential specifications in the input file.

10.2 Circuit Optimization

The generated circuits are basic representations. For production use, consider:

- Gate optimization using framework-specific tools
- Circuit compilation for specific hardware
- Noise-aware compilation for NISQ devices

10.3 Performance Considerations

- **Graph Size:** Performance degrades with graph size, especially for dense graphs
- **Clique Count:** Graphs with many maximal cliques take longer to process
- **Output Format:** Some frameworks generate more verbose output than others

11 Best Practices

1. **Start Small:** Test with small graphs before scaling up
2. **Verify Output:** Always check generated circuits for correctness
3. **Use Version Control:** Keep input files and generated outputs in version control
4. **Document Models:** Add comments to input files describing the model structure
5. **Test Multiple Frameworks:** Compare outputs across different frameworks
6. **Optimize Circuits:** Use framework-specific optimization tools after generation

12 Contributing

The MRF Compiler is an open-source project. Contributions are welcome! Areas for contribution include:

- Additional framework support
- Performance optimizations
- Extended input format support
- Documentation improvements
- Bug fixes and testing

13 License

Copyright (C) 2025, Shyamal Suhana Chandra

This software is provided as-is for educational and research purposes.

14 References

- Qiskit Documentation: <https://qiskit.org/documentation/>
- Cirq Documentation: <https://quantumai.google/cirq>
- PennyLane Documentation: <https://pennylane.ai/>
- Q# Documentation: <https://docs.microsoft.com/azure/quantum/>
- AWS Braket Documentation: <https://docs.aws.amazon.com/braket/>